```
> ourmodel <- lm(height ~ gender, data = genderheightdata)
> summary(ourmodel)

Call:
lm(formula = height ~ gender, data = genderheightdata)

Residuals:
   Min     1Q Median     3Q    Max
-7.500 -3.125  0.000  3.125  7.500

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  165.000      2.700  61.104 1.29e-09 ***
gendermale    12.500      3.819   3.273    0.017 *
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 5.401 on 6 degrees of freedom
Multiple R-squared:  0.641,  Adjusted R-squared:  0.5812
F-statistic: 10.71 on 1 and 6 DF,  p-value: 0.01696
```

We have a significant predictor (Gender) and model (indicated by the F-ratio).

- For a model with one predictor, the p values associated with the model (i.e., the F value) and the predictor are the same. For models with more than one predictor, this won't be the case.

- The Intercept coefficient (165) corresponds to the mean Height of our reference category (Female). The estimate gendermale (12.5) is the *difference* between our reference category and our Males. Females were taken as the reference category (i.e., the intercept) simply because R chooses this on an alphabetical basis (and *Female* comes before *Male*).