

- Our predictor doesn't have to be categorical though. We're using the `ageheightdata` here.

```
> ageheightdata
# A tibble: 8 x 3
  subject    age height
  <dbl>   <dbl>   <dbl>
1       1     22    170
2       2     21    180
3       3     19    175
4       4     23    185
5       5     15    160
6       6     17    170
7       7     16    165
8       8     17    165
```

Is Height predicted by Age?

```

> ourmodel <- lm(height ~ age, data = ageheightdata)
> summary(ourmodel)

Call:
lm(formula = height ~ age, data = ageheightdata)

Residuals:
    Min       1Q   Median       3Q      Max
-9.045 -2.104  1.646  3.201  3.557

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  126.281     11.411   11.067 3.24e-05 ***
age           2.398       0.602    3.984 0.00725 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.721 on 6 degrees of freedom
Multiple R-squared:  0.7257,    Adjusted R-squared:  0.6799
F-statistic: 15.87 on 1 and 6 DF,  p-value: 0.007252

```

For every increase in Age by 1, Height increases by 2.398. But of course, we know this relationship breaks down at a certain age - but for the data we have, we can fit a linear function.