

```
text_waroftheworlds <- books %>%  
  filter(title == "The War of the Worlds") %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words)  
  
text_underthesea <- books %>%  
  filter(title == "Twenty Thousand Leagues under the Sea") %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words)
```

- I then filter by book title and ‘unnest’ the text of each book so that we have one column in each of two dataframes corresponding to the words in each book...

```

> text_waroftheworlds
# A tibble: 60,513 x 3
  gutenber_id title word
      <int> <chr> <chr>
1         36 The War of the Worlds the
2         36 The War of the Worlds war
3         36 The War of the Worlds of
4         36 The War of the Worlds the
5         36 The War of the Worlds worlds
6         36 The War of the Worlds by
7         36 The War of the Worlds h
8         36 The War of the Worlds g
9         36 The War of the Worlds wells
10        36 The War of the Worlds 1898
# ... with 60,503 more rows

```

- If I hadn't use the `anti_join(stopwords)` call I would end up with a data frame like the above containing lots of common function words.